

Q1. What is the relationship between threads and processes?

Process:

A process is essentially a running instance of a program. It serves as a container for threads and possesses its own memory space to keep its code, data, and system resources isolated from other processes.

Memory space refers to the allocated memory that a process has for storing its code, data, and variables. Each process has its own separate memory space, ensuring isolation and preventing direct interference from other processes.

Threads:

Threads are the units of execution within a process. They can execute code independently while sharing the memory space and resources of the process they belong to. A process can have one or more threads.

File descriptors are system resource handles used by a process to interact with external resources such as files, sockets, etc. All threads within a process share access to these file descriptors, enabling them to read from or write to the same files or network connections.

A register set refers to the collection of registers a CPU has. Registers are small, fast storage locations within the CPU used to hold data temporarily during execution. Each thread has its own register set to keep track of its current state and to perform operations.

The stack is a region of memory where a thread keeps track of its function call history, local variables, and performs argument passing between functions. Each thread has its own stack to manage its execution flow.

A program counter (PC), also known as an instruction pointer, is a register that holds the memory address of the next instruction to be executed by a thread. Each thread has its own program counter to keep track of its position in the code.

Scheduling attributes determine how the operating system's scheduler manages the execution of threads. These attributes can include priority levels, scheduling policies, etc., and they are maintained on a per-thread basis, allowing the OS to decide which thread to run next based on these attributes.

Q2. Discuss advantages and disadvantages of supporting multi-threaded applications with (kernel-level) threads.

Advantages of Kernel-level Threads:

No Need for Yields in User Programs:

In user-level threading, programmers might need to insert yield operations to give other threads a chance to run. With kernel-level threads, the operating system's scheduler handles thread scheduling, so programmers don't need to explicitly yield control.

Thread.yield() in Java:

In Java, Thread.yield() is a method used to move a thread back to the ready state, allowing other threads a chance to execute.

```
public class YieldExample implements Runnable {
    public void run() {
        for (int i = 0; i < 5; i++) {
            System.out.println(Thread.currentThread().getId() + " Value " + i);
            Thread.yield();
        }
    }
}
```

Fair Amount of Execution Time:

The operating system's scheduler is designed to allocate CPU time fairly among threads, ensuring each thread gets a chance to execute.

Handling Blocking I/O:

If a thread performs a blocking I/O operation, the kernel can schedule another thread to run while the first thread is blocked. This helps to keep the CPU busy and improves the performance and responsiveness of the application.

Low Granularity Multitasking and Illusion of Parallelism:

Kernel-level threads provide a lower level of granularity in multitasking, allowing multiple threads within a process to execute seemingly in parallel, enhancing the performance and responsiveness of applications.

Multiprocessor Utilization:

Kernel-level threads can be scheduled on multiple processors, allowing true parallelism and better utilization of multi-core systems.

Disadvantages of Kernel-level Threads:**Need for Locks and Mutexes:**

Since the OS controls scheduling, threads can be preempted at any point, requiring careful synchronization (using locks and mutexes) around critical sections to prevent race conditions and data corruption.

Less Portability:

Kernel-level threading requires operating system support, which may vary across different systems, making the code less portable.

No Control over Scheduling Algorithm:

The scheduling algorithm used by the kernel may not always be suitable for the application's needs, and there's no way to replace or customize it without changing the kernel itself.

More Expensive Thread Management:

Creating, deleting, and managing kernel-level threads is more expensive compared to user-level threads as these operations require system calls which are more time-consuming.

Q3. Is putting security checks in the C library a good or a bad idea? Why?

Security Checks in C Library:**Voluntary Usage:**

The use of the C library is voluntary for user space programs. A program can choose to use or not use the C library functions, and can even bypass the library entirely to make system calls directly to the kernel.

Potential Bypass:

Even if security checks are embedded in the C library, a malicious program can bypass these checks by avoiding the C library and making system calls directly with malicious parameters.

Security Checks in Kernel Space:**Mandatory Path:**

All system calls, whether invoked directly or through the C library, must go through the kernel. This makes the kernel a mandatory path for programs seeking to access system resources.

Controlled Environment:

The kernel operates in a controlled, privileged environment (kernel space) where user space programs have no direct influence. This control allows for a secure environment to enforce security checks.

Cannot be Bypassed:

Since all access to system resources has to go through the kernel, security checks performed in kernel space cannot be bypassed by user space programs.

Example: Buffer Overflow Prevention:

Suppose the C library provides a safe string copying function called `safe_strcpy()` that includes a security check to prevent buffer overflow by ensuring the destination buffer is large enough for the source string plus the null terminator.

// Hypothetical `safe_strcpy` function in C library

```
void safe_strcpy(char *dest, const char *src, size_t dest_size) {  
    if (strlen(src) + 1 > dest_size) {  
        // Security check fails, handle error  
        fprintf(stderr, "Buffer overflow prevented\n");  
        exit(EXIT_FAILURE);  
    }  
    strcpy(dest, src); // Now, it's safe to copy  
}
```

Now, developers are encouraged to use `safe_strcpy()` instead of the standard `strcpy()` to prevent buffer overflows.

Bypassing the Library Security Check

Despite the availability of `safe_strcpy()`, nothing prevents a programmer (or a malicious actor) from using the standard `strcpy()` function or even making a system call directly to bypass the C library altogether.

```
char buffer[10];  
strcpy(buffer, "This string is too long and will cause a buffer overflow");
```

In this example, a buffer overflow occurs because `strcpy()` does not check whether the destination buffer is large enough to hold the source string.

Direct System Call

Moreover, a program could bypass the C library entirely and make direct system calls, which is a more advanced but entirely possible approach to interacting with the operating system:

```
syscall(SYS_write, file_descriptor, very_long_string, strlen(very_long_string));
```

In this snippet, `syscall()` is used to write data directly to a file, bypassing any security checks that might have been implemented in the C library's file writing functions.

Q4. What is a race condition? Give an example.

Explanation of Race Condition:

Uncoordinated Concurrent Access:

Race conditions arise when there is concurrent access to shared resources, and this access is not properly coordinated or synchronized.

Shared Resources:

The resources could be anything that multiple processes or threads use or depend on, such as shared memory, file systems, database records, or device registers.

The outcome of a race condition can be:

Incorrect behavior: The program does not work as expected.

Deadlocks: Two or more operations waiting for each other to complete, creating a cycle of dependencies.

Lost work: Updates from one operation may be overwritten by others, leading to lost or inconsistent data.

Example:

Initial state: Counter = 1

Process A reads counter: Counter = 1 (A's view: 1)

Process B reads counter: Counter = 1 (B's view: 1)

Process B updates counter: Counter = 2 (B's view: 2)

Process A updates counter (based on its earlier read): Counter = 2 (A's view: 2)

The crux of the problem is the lack of synchronization between Process A and Process B. If there had been a mechanism to ensure that reading and updating the counter were atomic operations (i.e., indivisible and uninterruptible), or if there had been a locking mechanism to ensure exclusive access to the counter, the race condition could have been avoided, and the counter would have been correctly incremented to 3.

Q5. What must the banker's algorithm know a priori in order to prevent deadlock?

The Banker's algorithm is a resource allocation and deadlock avoidance algorithm that tests for safety by simulating the allocation of predetermined maximum possible amounts of all resources, and then makes an "s-state" check to test for possible deadlock conditions for all other pending activities, before deciding whether allocation should be allowed to continue.

Maximal Resource Requirements:

The key piece of information required by the Banker's algorithm is the maximum resource requirements for each process in the system. Each process must declare the maximum number of each resource type that it may need.

Resource Tracking:

The algorithm assumes that it can track the availability and usage of each resource in the system. This includes knowing which resources are currently allocated to which processes, and which resources are still available for allocation.

Process Execution Simulation:

The Banker's algorithm simulates the execution of each process, assuming that processes receive the resources they require when needed, proceed to execute, and eventually finish and release all allocated resources back to the system.

Working of the Banker's Algorithm:**Initialization:**

Initially, the maximum demand of each process and the total number of available instances of each resource type are input to the system.

Resource Request:

When a process requests resources, the algorithm tentatively allocates the requested resources to that process and then checks whether the system would remain in a safe state if the allocation were allowed.

Safety Check:

A "safe" state is one where there exists at least one sequence of all process executions that does not result in deadlock. The algorithm simulates the execution of processes, checking whether each process can finish executing with the available resources.

Allocation Decision:

If the system remains in a safe state after the tentative allocation, the allocation is allowed to proceed. If not, the allocation is denied, and the process must wait for resources to become available.

Resource Release:

When a process completes its execution, it releases all its allocated resources, which then become available for other processes.

Q6. Describe the general strategy behind deadlock prevention and give an example of a practical deadlock prevention method.

Deadlock prevention is about structurally negating at least one of the four conditions required for a deadlock to occur. These four conditions are:

Mutual Exclusion:

Some resources are non-shareable; only one process can use them at a time.

Hold and Wait:

A process holds allocated resources while waiting for other resources.

No Preemption:

Resources cannot be forcibly taken away from a process; they must be released voluntarily.

Circular Wait:

A cycle of dependencies exists where processes are waiting for resources held by other processes in a circular chain.

A practical method to prevent deadlocks is to impose an ordering on resource allocation and ensure that all processes follow this ordering when requesting resources.

Suppose we have two resources: Resource A and Resource B. We can prevent deadlocks by establishing a rule that processes must always request Resource A before requesting Resource B.

Without Ordered Allocation:

Process 1 requests Resource B and holds it.

Process 2 requests Resource A and holds it.

Now, Process 1 is waiting for Resource A, and Process 2 is waiting for Resource B, forming a circular wait and causing a deadlock.

With Ordered Allocation:

All processes must request Resource A before requesting Resource B.

This order prevents the circular wait condition, and thereby, the deadlock.

Deadlock Prevention: This strategy negates at least one of the four necessary deadlock conditions, thereby structurally preventing the possibility of deadlocks.

Deadlock Avoidance: This strategy, exemplified by the Banker's algorithm, allows the system to enter potentially unsafe states but avoids deadlock by ensuring that resource allocation does not lead to an unsafe state.

Q7. Consider implementing a device interface (i.e., handling communication between CPU and a device) in a device controller rather than in the OS kernel. Which of the following statements is/are INCORRECT?

- A. Performance can be improved by hard-coded algorithms and utilising dedicated hardware.
- B. Device controller can introduce additional data buffering.
- C. The kernel is simplified by moving algorithms out of it.
- D. Improving algorithms requires a hardware update rather than just a device driver update.
- E. Bugs are less likely to cause an OS crash, and bugs are easier to fix.

Statement A (Performance Improvement):

Correct: By implementing a device interface in a device controller and utilizing dedicated hardware, it's possible to achieve better performance through hard-coded algorithms optimized for that particular hardware.

Statement B (Additional Data Buffering):

Correct: A device controller can introduce additional data buffering which may improve data handling efficiency between the device and the CPU.

Statement C (Kernel Simplification):

Correct: By moving algorithms out of the kernel and into a device controller, the kernel's complexity can be reduced, which might make the kernel more stable and easier to maintain.

Statement D (Algorithm Improvement):

Correct: If the algorithms are hard-coded into the hardware of the device controller, then improving or changing these algorithms would indeed require a hardware update, as

opposed to a software update (e.g., updating a device driver) if the algorithms were implemented in the OS kernel.

Statement E (Bug Handling):

Partly Incorrect:

Bug Impact: It's true that bugs in a device controller are less likely to cause an OS crash compared to bugs in the kernel since the kernel is a more critical part of the system.

Bug Fixing: However, bugs in a device controller with hard-coded algorithms could be harder to fix if they are embedded in hardware, as this could require a hardware update. In contrast, bugs in software (e.g., a device driver in the kernel) can be fixed with a software update, which is generally easier and faster to deploy.

Q8. Which statement about direct memory access (DMA) is CORRECT?

A. The DMA controller operates the memory bus by placing addresses on the bus to perform transfers with the help of the main CPU.

Correction: The DMA controller operates the memory bus directly without the help of the main CPU. It places addresses on the bus to perform transfers autonomously, which is a key feature of DMA as it frees up the CPU to do other tasks.

B. To initiate a DMA transfer, the host reads a DMA command block from the memory.

Correction: To initiate a DMA transfer, the host writes a DMA command block into memory. This block contains the necessary information for the transfer, and it is read by the DMA controller which then carries out the transfer.

C. DMA increases system concurrency by allowing executing instructions in parallel for a larger number of processes.

Correction: DMA does not execute instructions of processes. It merely handles data transfers between memory and I/O devices, freeing up the CPU. While this may indirectly allow for more concurrent execution of processes by the CPU, DMA itself does not execute process instructions.

D. In order to use DMA, hardware design becomes more complicated because the system must allow the DMA controller to be a bus master.

Correct Statement: The statement correctly points out that utilizing DMA introduces more complexity into the hardware design. For DMA to work, the system must allow the DMA controller to act as a bus master, meaning it can control the memory bus to manage data transfers. This is the correct statement about DMA.

E. All the above statements are correct.